

Chronosphere User Guide

Downloaded from Epic Games Confluence

Date: 2025-07-12 04:08:44

Original URL: <https://confluence-epicgames.atlassian.net/wiki/spaces/CDE/pages/81068700>

Document Level Classification

200

- [Introduction](#)
- [Finding Your Metrics](#)
 - [PromQL](#)
- [Tenants and Quotas](#)
 - [Quota Enforcement](#)
- [Advanced Features](#)
 - [Live Telemetry Analyzer](#)
 - [Telemetry Usage Analyzer](#)
 - [Metrics Aggregations](#)
 - [Drop Rules](#)
- [FAQ](#)
 - [How do quotas work during large events?](#)
 - [How do I see my quota limit and utilization?](#)
 - [How do I request a quota increase?](#)
 - [My quota is \\$X DPPS, what is this in real dollar terms?](#)
- [Glossary](#)

Introduction

Chronosphere is a new metrics storage system used by Epic. It is intended to be a replacement for New Relic Metrics use cases for APM, Infrastructure, Prometheus, and Custom metrics. You can think of it as a very large, scalable Prometheus server. It ingests metrics from our OTEL collectors and provides a PromQL interface to those metrics. We expose this as Prometheus datasources in Grafana, so you never need to interact with Chronosphere directly to make queries.

Chronosphere has a few key advantages over New Relic, and we believe you will have a better overall experience using it for metrics:

- **Significantly** faster query times. In an apples:apples test for the same data for the same app, Chronosphere was 4x to 12x faster than New Relic for loading individual panels in Grafana.
- Accuracy. Chronosphere is just a Prometheus implementation (M3) and implements the Prometheus spec. PromQL queries in New Relic are translated to NRQL, which can lead to inaccurate results, in particular for rate queries. Direct NRQL queries can have strange artifacts like random "spikes" or missing data points.
- Cost savings. Chronosphere has a lower overall cost and can help reduce costs further by helping reduce large metric footprints. The Data Platform team reduced 80% of its observability costs (around \$550k/year) through a combination of migrating to Chronosphere and using Chronosphere's governance and aggregation tools to reduce its footprint.

As of June 2024, by default, the OTEL pipeline in **all substrate clusters** will dual write to New Relic and Chronosphere. Any [metrics you send](#) to the cluster-local OTLP collector, or export with Prometheus, will be delivered to both vendors

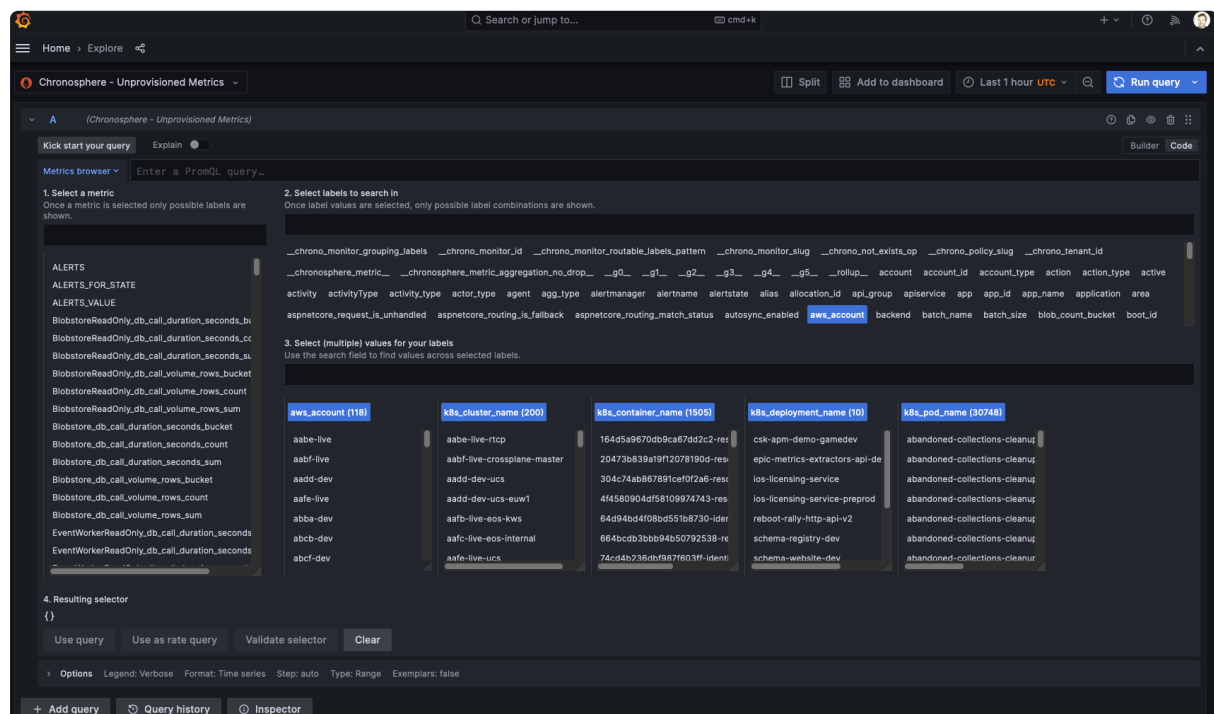
Finding Your Metrics

Chronosphere metrics are available in Grafana via Prometheus datasources starting with "Chronosphere". These are available in both

[Grafana Core](#) and [Grafana EOS](#). There are two categories of Chronosphere datasources:

- "Chronosphere Tenant" - this is datasource is for metrics with specific capacity reserved
- "Chronosphere - Unprovisioned Metrics" - this is where everything else goes. If our account capacity limits are reached, these metrics are at risk of being dropped. Before using any metrics for production use cases, contact [#ct-obs-support-ext](#) to have your metrics moved to a tenant.

If you're unsure of where to get started, begin with the Grafana Explorer UI. This will show you all metrics currently being sent to a specific Chronosphere tenant.



PromQL

If you are new to PromQL, have a look at the [Prometheus Query Basics](#) documentation. There is a bit of a learning curve if you are coming from NRQL, but Prometheus is fast which makes it fairly easy to iterate and learn. Here are some practical examples of queries that work in NRQL, translated to PromQL.

Computing Request Rate by Route

NRQL

```
SELECT rate(count(%.server%duration), 1 second) as `req/s` FROM Metric
```

PromQL

```
(sum(rate(http_server_duration_milliseconds_count{service_name="$service_name"}[1m])))
```

P99 Performance by Route

NRQL

```
SELECT percentile(%.server%duration, 0.99) FROM Metric WHERE service.name = '$service_name'
```

PromQL

```
histogram_quantile(0.99, sum by (le, http_method, http_route) (rate(http_server_duration_milliseconds_count[1m])))
```

Tenants and Quotas

Chronosphere capacity is provisioned up-front. This is by design to discourage run-away growth in metrics that lead to high observability costs. Our account capacity is provisioned into multiple metrics pools. The relationship between metric sources and the pool they go into is something we call a tenant. Tenants have a metrics datapoint-per-second (DPPS) quota, and all metrics going to a the same tenant are grouped under the same datasource. To illustrate with a practical example, let's look at how the Data org tenants work:

data-org-internal

- sources: fdbe-dev, cfef-live, edec-dev, fbfb-live (these are the AWS accounts emitting the metrics)
- quota: 40k DPPS
- grafana instance: grafana core
- datasource name: Chronosphere Tenant - data-org-internal

data-org-eos

- sources: eded-dev, dadb-live
- quota: 20k DPPS
- grafana instance: grafana eos
- datasource name: Chronosphere Tenant - data-org-eos

By default, your metrics go to a default pool that is de-prioritized and at risk at being dropped if account limits are reached. Before using Chronosphere metrics for production, reach out to [#ct-obs-support-ext](#) to have a tenant configured. The observability team will work with you to determine which tenant your metrics should go into, or whether a new one will be created. Typically new tenants will be created based on either organizational boundaries (Fortnite, Ecosec, Data, Mediatonic, etc) or when explicit separation is required (EOS for example). This mapping is done "on the fly" using dynamic labels and be reconfigured at any time, so it's possible to re-define what goes into a tenant in the future without losing any previous data. In other words, if we get it wrong, or there's a re-org in the future, we can change this later.

Quota Enforcement

Quotas are enforced when the account exceeds its DPPS capacity. As of June 2024, that capacity is 225,000 DPPS. Enforcement works like this:

- Once account capacity is reached, metrics in the default "unprovisioned" pool are dropped first. This is why you should have your metrics put in a tenant before relying on them for production.
- If capacity is still exceeded, tenants that are in violation of their quota are affected next
 - EKS dev cluster metrics are dropped first

- all other metrics not associated with an EKS live cluster are dropped after that
- if still in violation of the quota at this point, live cluster metrics will be dropped
- Tenants not in violation of their quota will not be affected by any enforcement mechanisms.

Advanced Features

Chronosphere has some advanced functionality for teams that want to get their metric costs under control. These are features only accessible in the Chronosphere UI. Since we think the typical user experience should start and end in Grafana, we don't grant this access by default, but if you need to use them please reach out. These are available as "power user" features.

Live Telemetry Analyzer

This allows us to peek at the "firehose" of metrics coming in and see large sources of writes, high cardinality metrics, etc. It's the first tool we use when asking "who or what is sending so much data?"

Live Telemetry Analyzer

Metrics

Traces

Capture live data

GD

Data phase

Accepted for storage

Pool

All

Priority

All

Add label filter

Labels

Showing first 500 results

Select labels to add to the label values table

Search

Label Keys	Unique Values	Appears In
<input checked="" type="checkbox"/> <code>__name__</code>	6552	100%
<input type="checkbox"/> <code>instance</code>	17788	100%
<input type="checkbox"/> <code>aws_account</code>	130	99%
<input type="checkbox"/> <code>epicotel_source</code>	5	99%
<input type="checkbox"/> <code>k8s_cluster_name</code>	196	96%
<input type="checkbox"/> <code>epic_source</code>	143	96%
<input type="checkbox"/> <code>k8s_node_name</code>	2960	88%
<input type="checkbox"/> <code>job</code>	50	83%
<input type="checkbox"/> <code>k8s_container_name</code>	1075	68%
<input type="checkbox"/> <code>k8s_namespace_name</code>	567	63%
<input type="checkbox"/> <code>k8s_pod_name</code>	15984	62%
<input type="checkbox"/> <code>namespace</code>	84	62%
<input type="checkbox"/> <code>service</code>	264	46%
<input type="checkbox"/> <code>pod</code>	6038	46%
<input type="checkbox"/> <code>endpoint</code>	409	41%
<input type="checkbox"/> <code>clusterName</code>	163	39%
<input type="checkbox"/> <code>cluster</code>	9	37%
<input type="checkbox"/> <code>container</code>	63	36%
<input type="checkbox"/> <code>nodeName</code>	1019	31%

Label values

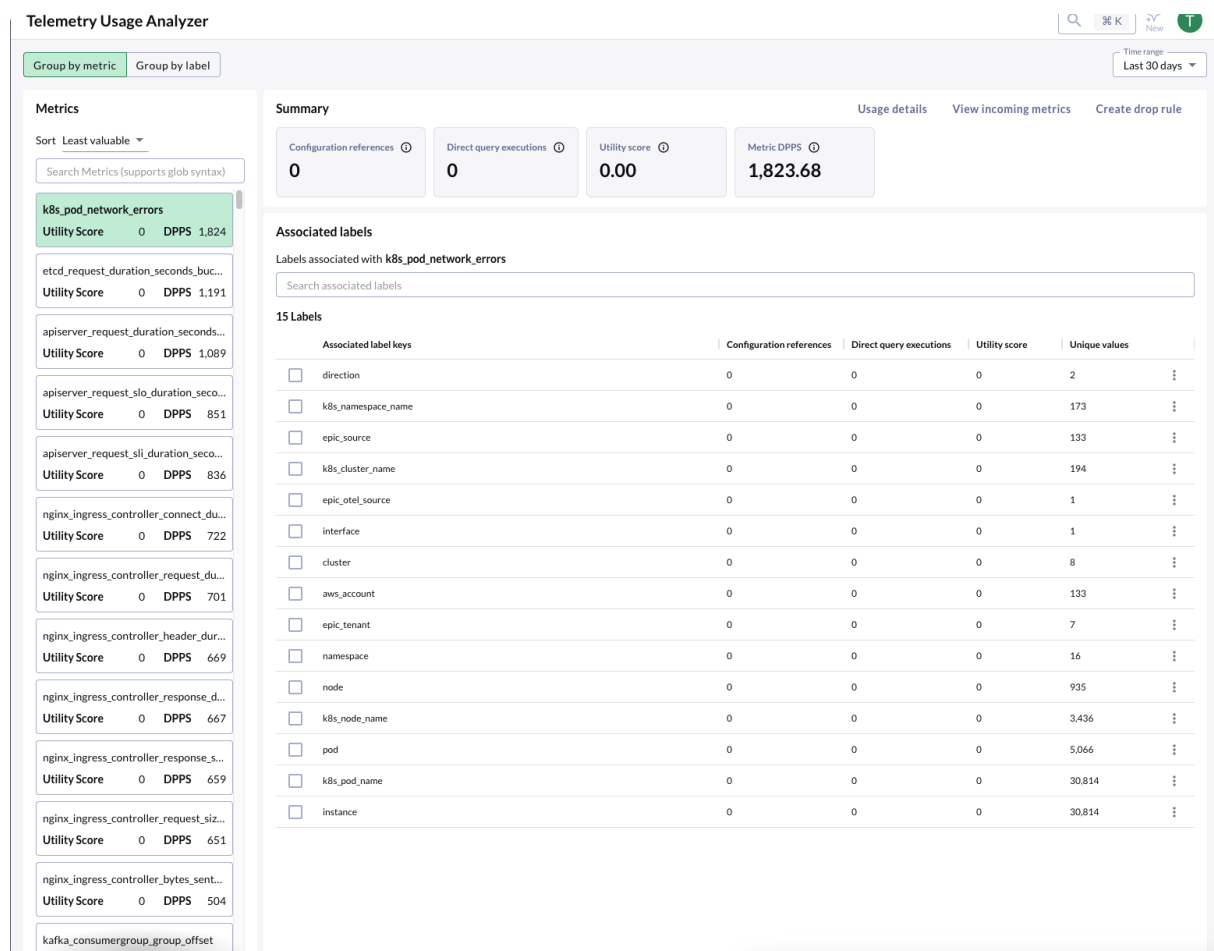
Showing first 500 results

Select label values to add to the label filter

<code>__name__</code>	<code>__metric_type__</code>	Avg. DPPS	Current DPPS
<code>kube_pod_status_reason</code>	-not set-	1659	2053
<code>kube_pod_status_phase</code>	-not set-	1659	2075
<code>k8s_pod_network_io</code>	-not set-	1597	2136
<code>k8s_pod_network_errors</code>	-not set-	1597	2136
<code>kube_node_status_condition</code>	-not set-	1399	1020
<code>container_memory_working_set</code>	-not set-	1191	1615
<code>container_cpu_utilization</code>	-not set-	1191	1615
<code>container_memory_usage</code>	-not set-	1189	1613
<code>apiserver_request_duration_seconds_bucket</code>	-not set-	1082	0
<code>kube_pod_status_ready</code>	-not set-	995	1245
<code>etcd_request_duration_seconds_bucket</code>	-not set-	920	0
<code>nginx_ingress_controller_request_duration_seconds</code>	-not set-	907	0
<code>loki_request_duration_seconds_bucket</code>	-not set-	896	1500
<code>nginx_ingress_controller_response_size_bucket</code>	-not set-	887	0
<code>nginx_ingress_controller_response_duration_second</code>	-not set-	853	0
<code>kube_pod_container_resource_requests</code>	-not set-	830	1316
<code>metric_enveloping_timer_seconds_bucket</code>	-not set-	802	408
<code>k8s_pod_uptime</code>	-not set-	798	1068
<code>k8s_pod_memory_working_set</code>	-not set-	798	1068
<code>k8s_pod_memory_usage</code>	-not set-	798	1068

Telemetry Usage Analyzer

This tool helps us see which metrics are being used. It takes a lot of guesswork out of "do I even need this?"



Metrics Aggregations

Metrics aggregations can be used for a number of purposes but a major one is to reduce high cardinality metrics and high DPPS by dropping unneeded dimensions. This is very useful, for example, when you have no control over the exported metrics and it's sending high cardinality labels that you don't want. To illustrate the concept, imagine you have some metrics coming in like this:

```
http_request_count{service="web", k8s_pod_name="foo", http_status_code=200}
http_request_count{service="web", k8s_pod_name="bar", http_status_code=200}
http_request_count{service="web", k8s_pod_name="xyz", http_status_code=200}
http_request_count{service="web", k8s_pod_name="foo", http_status_code=201}
http_request_count{service="web", k8s_pod_name="bar", http_status_code=201}
http_request_count{service="web", k8s_pod_name="xyz", http_status_code=201}
```


You have a lot of pods emitting these metrics but you don't care about this metric at the pod level. Rather than store one metric per pod, you can use an aggregation rule to sum up these metrics and drop the pod name. This would produce a metric like:

```
http_request_count{service="web", http_status_code="200"} 4
http_request_count{service="web", http_status_code="204"} 12
```

Processing metrics like this is not free but it is substantially cheaper (about 1/3rd the price) of ingesting the raw metrics. Aggregation rules are managed in [this terraform repository](#). If you would like help setting aggregation rules, please reach out to [#ct-obs-support-ext](#).

Drop Rules

Drop rules let you drop noisy metrics in a pinch at the ingestion layer at Chronosphere. While it's more efficient to simply not emit the metric in the first place, the ability to drop metrics without changing any code or re-deploying is useful especially in an emergency. Drop rules are managed in [this terraform repository](#). If you would like help setting drop rules, please reach out to [#ct-obs-support-ext](#).

FAQ

How do quotas work during large events?

Quotas are not enforced if the account-wide limit has not been reached. Ahead of a large-scale event, the observability team will request additional capacity from Chronosphere, which will effectively raise any quota you currently have. **tl;dr** you do not need to worry about micromanaging your quotas ahead of an event.

How do I see my quota limit and utilization?

See the Chronosphere Quota Utilization dashboard [here](#). For EOS grafana, the dashboard is [here](#). We do not yet have limit values in these dashboards, so to inspect your limit, see [this terraform file for now](#).

How do I request a quota increase?

Just contact [#ct-obs-support-ext](#). We can work with you to raise your quota.

My quota is \$X DPPS, what is this in real dollar terms?

This will likely change in the future as we get better bulk pricing, but currently each 10000 DPPS allocated to your tenant quota runs approximately \$15000/year.

Glossary

Tenant - a logical grouping of related metric sources that roll up under the same quota and are accessed through the same grafana datasource. These are usually aligned with a particular business unit or organization.

DPPS - datapoints per second, or the number of metrics being ingested per second

Page Information:

Page ID: 81068700

Space: Cloud Developer Platform

Downloaded: 2025-07-12 04:08:44