

# Custom Kubernetes Node Types for Karpenter Autoscaling

---

Downloaded from Epic Games Confluence

Date: 2025-07-12 04:09:10

Original URL: <https://confluence-epicgames.atlassian.net/wiki/spaces/CDE/pages/81068359>

Document Level Classification

300

- [Intended audience: Development Teams and DevOps Engineers](#)
- [Making Changes to Karpenter Nodepools](#)
- [Custom Nodepool Creation](#)
  - [Dedicated ARM64 Example](#)
  - [NVIDIA GPU Example](#)
- [Assigning Workloads to Your Tainted Nodes](#)
- [Deploying Your Nodepool](#)

## Intended audience: Development Teams and DevOps Engineers

Kubernetes cluster autoscaling is handled by [Karpenter](#). Deployments to our clusters are managed by our [ArgoCD](#) repository, but only Substrate engineers have access to approve and push out changes (they could effect every cluster in the company if done wrong). If you would like a new pool of nodes attached to your cluster of a specific node type or architecture, this guide will walk you through that process.

## Making Changes to Karpenter Nodepools

Before making changes to how autoscaling works in your cluster, you need to know about the Karpenter [CRD](#) used to configure it. This custom resource used to define autoscaling of nodes by Karpenter is called 'nodepool' and the documentation for it is here: <https://karpenter.sh/docs/concepts/nodepools/>

The overall configuration process goes in the following order, but you will only be making changes to the Wrapper Helm Chart with the intent that the Karpenter Nodepool CRDs in your cluster will be changed. If you want to see the current configuration(s) in your cluster, you can run `kubectl get nodepool -A -o yaml`.

```
ArgoCD Repository --> Wrapper Helm Chart (by Epic) --> Karpenter  
Helm Chart --> Karpenter NodePool CRD
```

## Custom Nodepool Creation

To configure a new nodepool, you will need to branch the [ArgoCD](#) repository. Once your branch has been created, find your cluster in the [cluster-configs](#) directory. This file controls the cluster-specific overrides values sent to all Helm charts that are sent to your cluster. In this file we can provision an additional nodepool for use with your cluster.

See the following example of a custom nodepool. You are able to nest any value [in this Karpenter wrapper chart](#) under the `karpenter:` section below. You can see the [values supported by the chart in the template here](#). In this example, the new nodepool will be named custom (``kubectl get nodepool custom``) and will use `amd64` architecture nodes in the `instanceType` values list.

If you wanted to use GPU accelerated nodes, you would simply change the node type values to [AWS GPU node types](#). If you wanted ARM nodes, you would change the architecture to `arm64`. To see a list of all things that can be filtered on, check [the supported labels here](#).

To see a list of the resources you can request by instanceType, go [here](#) and [here](#).

## Dedicated ARM64 Example

### dead-dev-dontwork.yaml

```
karpenter:  
  values:  
  
    nodepools:  
      custom:  
        enabled: true  
        taints:  
          - key: epicgames.com/my-special-taint  
            effect: NoExecute  
        labels:  
          use: dedicated  
          app: jenkins  
        requirements:  
          zone:  
            values: []  
          arch:  
            values:  
              - amd64  
          capacityType:  
            values:  
              - on-demand  
          instanceType:  
            values:  
              - c5.9xlarge  
              - c5.12xlarge  
              - m5a.12xlarge  
              - c6i.8xlarge  
              - c6i.12xlarge
```

```
- c6a.8xlarge
- c6a.12xlarge
- r6i.8xlarge
- r6i.12xlarge
- r5.8xlarge
- r5.12xlarge
- r5a.12xlarge
nodeClassRef:
  name: default
disruption:
  consolidationPolicy: WhenEmptyOrUnderutilized
```

## NVIDIA GPU Example

It is feasible to *not* include a `taints` section as is done in the example below. Given that Karpenter will provision nodes based on the lowest node pricing at the time of pods being Pending, this should work out fine for GPU nodes. However, if you have a mix of system architectures without taints or explicit node selectors on all your workloads, you will find that pod will schedule onto nodes of the wrong architecture and thereafter will fail to start. In short, you can leave taints off of nodepools that use GPU acceleration, but you should not remove the taints for nodepools that are of non-amd64 architecture or else your pods will unexpectedly schedule onto nodes of the wrong architecture.

### WARNING

It appears that currently instance types starting with p3. or p4. do not properly contain the "[nvidia.com/gpu](https://nvidia.com/gpu)" resource. If you want to use NVIDIA GPUs, please use the instance types beginning with g5.

### **dead-dev-dontwork.yaml**

```
# Remember to check that you have enough instance quota for GPU
nodes!
karpenter:
```

values:

nodepools:

```
enabled: true
labels:
  use: not-mining
taints:
  - key : epicgames.com/gpu
    effect: NoExecute
requirements:
  zone:
    values: []
  arch:
    values:
      - amd64
  capacityType:
    values:
      - on-demand
  instanceType:
    values: []
  instanceCategory:
    values:
      - g
  instanceGPUManufacturer:
    values:
      - nvidia
nodeClassRef:
  name: default
disruption:
  consolidationPolicy: WhenEmptyOrUnderutilized
```

You can see an example of a deployment that consumes GPUs [here](#).

## Assigning Workloads to Your Tainted Nodes

When a pod with a node selector is detected as Pending, Karpenter will automatically determine which nodes need to be provisioned to satisfy the demands of that pod. This means that your next step is to configure the workloads that you want to provision onto your new tainted Karpenter nodes. The [Assigning Workloads to Tainted Nodes](#) article has all the information you need to do that.

## Deploying Your Nodepool

Once your nodepool and workload specs are written, make a pull request to your cluster spec in the [substrate/argocd](#) repository. Hosting team may be able to help you out. A member with access to the ArgoCD UI will verify your pull request, merge it, and deploy your change to your cluster.

As soon as your new nodepool is in place, which you can check with `kubectl get nodepool`, you can safely merge your updated spec with tolerations and node label selectors. Now, when new pods are Pending, Karpenter will automatically provision tainted nodes for your workload that no other workloads will be able to use!

---

### Page Information:

Page ID: 81068359

Space: Cloud Developer Platform

Downloaded: 2025-07-12 04:09:10