

# Deploying your first application to Substrate

---

Downloaded from Epic Games Confluence

Date: 2025-07-12 04:06:55

Original URL: <https://confluence-epicgames.atlassian.net/wiki/spaces/CDE/pages/81068426>

Document Level Classification

200

- [Introduction](#)
  - [Prerequisites](#)
  - [Placeholders](#)
- [Tutorial](#)
  - [01 - Create a sample application](#)
    - [A - Initialize working directory and files](#)
    - [B - Build and Run locally](#)
    - [C - Publish to Artifactory](#)
  - [02 - Create a Helm Chart](#)
    - [A - epic-app](#)
  - [03 - Configure authentication and authorization](#)
    - [A - Add UAP](#)
    - [B - Configure UAP](#)

- [04 - Configure ingress](#)
  - [A - Create a Service](#)
  - [B - Configure Ingress and Authorize Ingress Traffic](#)
    - [Internal vs Internet Facing Load Balancers](#)
- [05 - Using secrets](#)
  - [Option 1A - Publish and share secrets with Substrate Vault](#)
  - [Option 1B - Access and use secrets in your application using ESO](#)
  - [Option 2A - Publish and share secrets with Substrate Vault and SSSM](#)
  - [Option 2B - Access and use secrets in your application using vault injector](#)
- [06 - Deploy the application](#)
- [07 - Verify deployment](#)
  - [A - Access the deployed application](#)
  - [B - View logs using kubectl](#)
  - [C - View logs using Grafana](#)
  - [D - View metrics using Grafana](#)

# Introduction

*This tutorial is designed to give you an introduction to Substrate using a sample application deployed to your Substrate infrastructure.*

## **Prerequisites**

1. Complete the [Prerequisites](#) to ensure you have the appropriate access and local developer toolchain setup.
2. Verify access to your [Substrate infrastructure](#). This tutorial interacts with your Substrate infrastructure using an [authenticated terminal session](#).

## Placeholders

The instructions and code snippets in this tutorial use the following placeholders. Make sure to replace them appropriately based on your Substrate infrastructure.

Placeholder	Description	Example
<SUBSTRATE_ACCOUNT>	The name of your Substrate account.	abcd-dev
<SUBSTRATE_PREFIX>	The prefix in your Substrate account name	abcd
<SUBSTRATE_CLUSTER>	The name of your Substrate cluster.	abcd-dev-substrate
<SUBSTRATE_NAMESPACE>	The name of your Substrate cluster namespace.	team-foobar
<DOCKER_REPO>	The name of your Artifactory docker repository.	foobar-docker-dev
<HELM_REPO>	The name of your Artifactory helm repository.	foobar-helm-dev
<USERNAME>	Your Okta username.	john.doe
<USERNAME_ALNUM>	An alpha-numeric version of your Okta username.	johndoe
<ARTIFACTORY_ACCESS_TOKEN>		abcdef123456

	The Identity Token associated with your <a href="#">Artifactory Profile</a> .	
--	---	--

# Tutorial

## 01 - Create a sample application

Substrate is optimized for running containerized workloads. Let's create a simple web application built using [nginx](#) that is packaged as a [container image](#).

### A - Initialize working directory and files

Create a directory for the application:

```
# Create a working directory for the application
mkdir ~/substrate-example
cd ~/substrate-example

# Initialize as a git repository
git init -b main .
```

We recommend that you use `git` to commit changes as you are creating and updating files. Doing so will make it easier to revert changes as well as use `git-diff` and/or `git-log` to review changes made as you progress through the steps in this tutorial.

Populate the application directory with the following files:

Filename	Description	Contents
----------	-------------	----------

<div data-bbox="213 1025 325 1061" data-label="Text"> <p>app.js</p> </div>	<div data-bbox="429 891 608 1189" data-label="Text"> <p><i>Simple application written in JavaScript and uses <a href="#">njs</a>.</i></p> </div>	<div data-bbox="708 215 1596 1787" data-label="Text"> <pre>function home(r) {   r.headersOut["Content-Type"] = "application/"   r.return(200, JSON.stringify({ message: "OK" })  function showHeaders(r) {   r.headersOut["Content-Type"] = "application/"   r.return(200, JSON.stringify({ headers: r.he })  function showEnvVars(r) {   r.headersOut["Content-Type"] = "application/"   r.return(200, JSON.stringify({ environmentVa })  function showAuth(r) {   r.headersOut["Content-Type"] = "application/"    function _jwtDecode(data) {     var parts = data       .split(".")       .slice(0, 2)       .map((v) =&gt; Buffer.from(v, "base64url").       .map(JSON.parse);     return { headers: parts[0], payload: parts    }    r.return(200, JSON.stringify(_jwtDecode(r.he })  export default { home, showHeaders, showEnvVar</pre> </div>
<div data-bbox="213 1966 373 2002" data-label="Text"> <p>app.envsh</p> </div>		

	<i>Application container entrypoint.</i>	<pre>#!/bin/sh  # Load any secrets available from vault if test -d "/vault/secrets"; then     find "/vault/secrets" \         -follow \         -type f \         -print \           sort -V \           while read -r f; do         . "\${f}"     done fi</pre>
nginx.conf	<i>nginx configuration file.</i>	<pre>load_module modules/nginx_http_js_module.so;  events { }  http {      js_path "/etc/nginx/njs/";     js_import app.js;      server {         listen 80;          location / {             js_content app.home;         }          location /health {             return 200;         }     } }</pre>

		<pre>     }      location /show-headers {         js_content app.showHeaders;     }      location /show-env {         js_content app.showEnvVars;     }      location /show-auth {         js_content app.showAuth;     } } </pre>
Dockerfile	<a href="#"><u>Instructions</u></a> <i>to build a container image.</i>	<pre> # The sample application uses the nginx _Base # Reference: https://gallery.ecr.aws/nginx/nginx:alpine FROM public.ecr.aws/nginx/nginx:alpine  # Add application files to the image COPY nginx.conf /etc/nginx/nginx.conf COPY app.js /etc/nginx/njs/app.js  # Configure entrypoint COPY app.envsh /docker-entrypoint.d/app.envsh RUN chmod +x /docker-entrypoint.d/app.envsh </pre>

## ***B - Build and Run locally***

*Build the container image using the following command:*

```
docker build -t substrate-example .
```

*Run the container locally to verify that it works:*

```
# Start the container, with the application listening on port 80 (you m
docker run --rm -p 80:80 --name substrate-example substrate-example

# Verify that the container is running
docker ps

# -----

# Use curl to invoke the application

# The / route should return a JSON response '{"message":"OK"}'
curl -LSsf -- http://localhost/ | jq -S .

# The /health route should return an empty 200 OK response
curl -LSsfv -- http://localhost/health

# The /show-headers route should respond with the request headers
curl -LSsf -H "x-foo: bar" -- http://localhost/show-headers | jq -S .

# The /show-env route should respond with the environment variables ava
curl -LSsf -- http://localhost/show-env | jq -S .

# The /show-auth route should respond with the payload of the Authoriza
# Use aop to generate a JWT
curl -LSsf -H "Authorization: Bearer $(aop uas)" -- http://localhost/sh

# -----

# To observe the logs generated by the application, use:
docker logs substrate-example
```



```
# Finally, stop the container
docker stop substrate-example
```

## **C - Publish to Artifactory**

*Use the following commands to publish the container image to Artifactory:*

```
# Authenticate to Artifactory
docker login -u '<USERNAME>' -p '<ARTIFACTORY_ACCESS_TOKEN>' artifacts.

# Tag the container image
# Remove the single quotes
docker tag substrate-example 'artifacts.ol.epicgames.net/<DOCKER_REPO>/

# Publish the container image
docker push 'artifacts.ol.epicgames.net/<DOCKER_REPO>/substrate-example
```

## **02 - Create a Helm Chart**

*In the previous section you packaged the sample application as a container image. To run the container image on Kubernetes, you need to also provide configurations to describe the runtime requirements. For example, how many instances of the application should be running? What are the minimum CPU and memory requirements? What ports need to be enabled? and many more. You can define these requirements within a [Helm Chart](#) - a [chart](#) is a collection of files that describe a related set of Kubernetes resources. Charts are packages that can be published, consumed and extended (similar to software packages and container images).*

### **A - epic-app**

[epic-app](#) is an opinionated Helm Chart offered to all teams across Epic to abstract away typical resources and configurations for workloads running on Substrate infrastructure. Create the following files in your application directory to define a new chart that depends on `epic-app` :

Filename	Description	Contents
deploy/chart/ Chart.yaml	A YAML file containing information about the chart.	<pre>##### # HELM CHART FOR `substrate-examp #####  # # Reference: # * https://helm.sh/docs/topics/ch #  apiVersion: v2 name: substrate-example version: "1.0.0-&lt;USERNAME_ALNUM&gt;" description: Substrate example app  # Add a dependency on epic-app (ht dependencies:   - name: epic-app     repository: @substr-helm       # epic-app versions and as     # https://github.ol.epicgames.     version: "&gt;= 2.4.6, &lt; 2.6.0"  #####</pre>
deploy/chart/ values.dev.yaml	The dev configuration values for this chart. You may eventually have a similar values.live.yaml for live.	<pre>##### # Configuration values for substra #####  # # NOTE:</pre>

```
# These are values passed to the e
# under a top level property that
# `Chart.yaml`. By convention, the
#
epic-app:
  #
  # Reference:
  # * https://github.ol.epicgames.
  #

resourceTags:
  service: "<name of the service"
  owner: "<name of the owner>"
  contact: "<email of the owner>"
  euid: "<EUID in Epic's service"

containers: # This is a map of c
# Application container
substrate-example:
  image:
    name: "artifacts.ol.epicga
    tag: "1.0.0-<USERNAME_ALNU
  ports:
    - 80
  environment:
    LOG_LEVEL: INFO
```

```
#####
```

*Initialize and validate your new chart using the commands below:*

```
# List available helm repositories
helm repo list
```

```
# If not listed above, add the "substr-helm" repository hosted on Artifactory
helm repo add substr-helm https://artifacts.ol.epicgames.net/artifactory/substr-helm

# get/update the epic-app dependency
helm dependency update deploy/chart
```

*Running the commands above will add a couple of new files in your application directory:*

<b>Filename</b>	<b>Description</b>
<code>deploy/chart/charts/epic-app-*.tgz</code>	<i>This is the epic-app chart downloaded to disk. The containing charts directory will contain downloaded copies of dependencies and should be ignored by version control (via <code>.gitignore</code> )</i>
<code>deploy/chart/requirements.lock</code>	<i>This is a dependency lock file that can be used to rebuild the dependencies to an exact version. It is recommended to commit this file to version control.</i>

## **03 - Configure authentication and authorization**

For applications that rely on <https://confluence-epicgames.atlassian.net/wiki/spaces/CE/pages/93488860> (UAS), authentication and authorization can be delegated to the <https://confluence-epicgames.atlassian.net/wiki/spaces/CE/pages/93488602> (UAP). UAP is available as a container image and can be deployed alongside your application (i.e., a sidecar-proxy to your upstream application container).

## A - Add UAP

UAS connection issues

If you get such errors "<https://uas.beef.live.use1a.on.epicgames.com/api/v1/keys>": dial tcp i/o timeout

*you should ask #cloud-ops-support-ext to authorize your account on UAS service.*

*To add UAP to your deployment, update the `containers` section of `values.dev.yaml` to include the following snippet:*

**Mosaic** macros cannot be exported to this format.

## B - Configure UAP

*UAP, by default, will ensure that all HTTP requests are authenticated and authorized before forwarding the request to your application. However, there are scenarios where unauthenticated requests must be allowed. For the sample application, requests to the /health route should not require authentication to allow for health-checks. UAP [allows for this scenario](#) using a configuration file. You can provide the configuration file to the UAP container using a combination of [ConfigMaps and Volumes](#).*

*To provide a custom route configuration file to UAP, replace or update the UAP container definition and include the additional configurations from the snippet below:*

**Mosaic** macros cannot be exported to this format.

## 04 - Configure ingress

*The application is a service and is intended to be invoked via HTTP requests. Enabling ingress requires multiple steps:*

1. Create a [Service](#) – In Kubernetes, a Service is an abstraction which defines a logical set of Pods and a policy by which to access them.

2. Configure [Ingress](#) – Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.
3. Control network traffic – Define rules that allow network traffic to reach your ingress.

## **A - Create a Service**

To expose your application as a Service, update `values.dev.yaml` to include the following snippet:

When using UAP, ingress must be configured to point to UAP and not the application. Otherwise, the application will receive requests that are not authenticated or authorized.

**Mosaic** macros cannot be exported to this format.

## **B - Configure Ingress and Authorize Ingress Traffic**

Next, you need to configure ingress. You can use an [Application Load Balancer](#) for this. In Kubernetes, an [Ingress controller](#) is responsible for fulfilling the ingress. Your Substrate cluster comes with the [AWS Load Balancer Controller](#) deployed and can be used to provision an Application Load Balancer.

The Application Load balancer will be configured to listen on port 80 (HTTP) and port 443 (HTTPS), with all traffic on port 80 being redirected (HTTP-301) to port 443. To use an HTTPS listener, you need a valid SSL certificate provisioned in [AWS Certificate Manager \(ACM\)](#). Follow [these instructions](#) to create a certificate in ACM.

To add an Application Load Balancer to your application, update `values.dev.yaml` to include the following snippet. Replace `<ACM_CERTIFICATE_ARN>` with the value for your certificate.

## Internal vs Internet Facing Load Balancers

The preferred default method for configuring your load balancer is `internal` so your service can be accessed internally via the [Service Network](#). You would only configure your load balancer as internet-facing if your service would be used to serve external customers. Please be sure to use the correct examples below. They are suffixed by either `- internal` or `- internet facing` depending on your use case.

*If you are using the default preferred method to load balance traffic to an **internal** Application Load Balancer the internal examples below allow traffic via the Service Network, Office, and VPN.*

*If you are using an **internet facing** Application Load Balancer, you will need to specify a prefix list, or inbound CIDR, or a Security Group annotation (you can only choose one method of annotation). In the case of using a prefix list or inbound CIDR ingress annotation, the load balancer controller manages the security group. I.E. the attached security group will include a combination of the listeners annotations and the addresses of the specified prefix list or inbound CIDR. In the case of using a prefix list or security group, these would need to exist in AWS in order to specify the name. Likely these would have been created using Terraform. For more information refer to [Managing inbound traffic to your application](#). For the internet facing example below we will use the prefix list annotation.*

**Mosaic** macros cannot be exported to this format.

*Traffic to the Application Load Balancer is controlled by rules associated with a [security group](#) that is attached to it. By default, a new security group is automatically created and attached to the Application Load Balancer.*

## 05 - Using secrets

Applications use sensitive configurations, or secrets, when accessing other backend or downstream resources. For example, a database password or an API key. For a Substrate application, secrets are stored and managed in <https://confluence-epicgames.atlassian.net/wiki/spaces/CE/pages/93487474>. Your application can then be configured to use these secrets.

Starting from August 2024, it is mandatory that all new Kubernetes clusters use External Secrets Operator (ESO). We strongly recommend using ESO, but the Vault Injector sidecar will still remain available with limited support, bug fixes will be provided if needed to maintain core functionality.

For the purposes of this tutorial we will be using Option 1 below to use ESO for secrets. Reference [Using External Secrets Operator \(ESO\) in epic-app to inject secrets](#) for documentation on using ESO.

**There are currently 2 ways to use secrets. Both options are presented below. You only have to choose 1 option.**

**Option 1:** Publishing secrets in Substrate Vault, then using External Secrets Operator in epic-app to consume those secrets.

**Option 2:** Publishing and sharing secrets with Substrate Vault and SSSM, then using vault injector in epic-app to consume those secrets

### **Option 1A - Publish and share secrets with Substrate Vault**

Lets assume that the sample application needs 2 secrets to access a database. This will be a username and password. You would first need to create the secrets in substrate vault. The instructions for creating a secret in Substrate Vault for use with ESO in epic app can be found in the



document called [Using External Secrets Operator \(ESO\) in epic-app to inject secrets](#).

### **Option 1B - Access and use secrets in your application using ESO**

For applications to access secrets, you can enable externalSecrets in epic-app. When the application is deployed, ESO will make the secrets available to the containers in the pod. In this example

### **Option 2A - Publish and share secrets with Substrate Vault and SSSM**

Let's assume that the sample application needs a secret named `foo` with a value `bar`. Use the [Substrate vault instructions](#) to publish the secret and share it with your Substrate cluster. The instructions below assumes that you have a secret published with a path `secret/path/to/foo`, make sure to update it to reflect the path used in Vault.

### **Option 2B - Access and use secrets in your application using vault injector**

For applications to access secrets, you can use the [HashiCorp Vault Agent Injector](#). The Vault Agent Injector is already installed in a Substrate cluster, you need to use annotations to enable it for your application. When the application is deployed, the Vault Agent Injector will fetch the secrets configured via annotations and make them available as files to your application. In this example, a file called `/vault/secrets/foo` will be available. The application container entrypoint (`app.envsh`) reads the file and passes on the secrets as environment variables to the application process.

To access secrets within the application, update `values.dev.yaml` to include the following snippet:

**Mosaic** macros cannot be exported to this format.

## 06 - Deploy the application

The application can be deployed using [helm install](#).

This requires an [authenticated terminal session](#). Make sure you have authenticated with the account you intend to deploy to as well as the correct namespace.

```
# Examine the manifests that epic-app and helm generate for the application
helm template 'substrate-example-<USERNAME_ALNUM>' ./deploy/chart -f ./values.yaml

# Deploy the application
helm install 'substrate-example-<USERNAME_ALNUM>' ./deploy/chart -f ./values.yaml -n '<SUBSTRATE_NAMESPACE>'

# Inspect the deployment
helm status 'substrate-example-<USERNAME_ALNUM>' --show-resources -n '<SUBSTRATE_NAMESPACE>'

# Review the list of resources created
helm get all 'substrate-example-<USERNAME_ALNUM>' -n '<SUBSTRATE_NAMESPACE>'
```

## 07 - Verify deployment

### A - Access the deployed application

For internal Service Network use cases browse to `https://substrate-example-`

`<USERNAME_ALNUM>.<SUBSTRATE_ACCOUNT>.dev.use1a.internal.epicgames.com`

For Internet facing use cases browse to `https://substrate-example-`

`<USERNAME_ALNUM>.<SUBSTRATE_ACCOUNT>.dev.use1a.internal.epicgames.com`

Review the responses from the following routes (they should be similar to [what was observed locally](#)):

- `/health`
- `/show-headers`
- `/show-env`
- `/show-auth`

## **B - View logs using `kubectl`**

To view logs from your application on the terminal you can use [kubectl logs](#) :

```
# Get the name of the pod
kubectl get pods -n '<SUBSTRATE_NAMESPACE>' -o jsonpath='{.items[0].met

# Get logs generated since the last hour for the application container
kubectl logs --since=1h -n '<SUBSTRATE_NAMESPACE>' -c substrate-example
```

## **C - View logs using Grafana**

Container logs are, by default, forwarded to Loki and visualized in Grafana. To view the logs in Grafana, you can:

1. Go to [Loki User Guide](#)
2. Follow the instructions under **Finding Your Logs** and **Querying Your Logs**

## **D - View metrics using Grafana**

Substrate cluster metrics, by default, are forwarded to Chronosphere and visualized in Grafana. To view available metrics for your application, you can:

1. Go to the [Chronosphere User Guide](#)
2. Follow the instructions under **Finding Your Metrics**

---

**Page Information:**

Page ID: 81068426

Space: Cloud Developer Platform

Downloaded: 2025-07-12 04:06:55