

How To Enable GPU Access for EKS Cluster

Downloaded from Epic Games Confluence

Date: 2025-07-12 04:07:13

Original URL: <https://confluence-epicgames.atlassian.net/wiki/spaces/CDE/pages/81068307>

Document Level Classification

[200](#)

- [Introduction](#)
- [Worker Nodes](#)
- [Karpenter Nodes](#)
- [Enabling GPU Support in Kubernetes](#)
 - [Deployment via helm](#)
- [Time-Slicing GPUs in Kubernetes](#)

Introduction

This guide will show how to enable GPU instances access from EKS Clusters to enable Graphics Accelerated Applications or Machine Learning projects. The following procedure can be made using two kinds of node management:

Worker Nodes

- Worker nodes require to perform a specific count of g5 nodes addition as worker nodes to your EKS Cluster, this is helpful if you already had dimensioned the count of replicas/pods or GPU you'll require along the service/project is running
- In this mode of worker node management, you add the instance type and a certain quantity of nodes directly to your eks cluster, this mode leaves you to add min, max, and desire count of nodes but autoscaling is not managed dynamically by Karpenter.
- In this mode you require to modify your terraform [substrate/infrastructure](#) for your specific EKS Cluster were you want to add these GPU Nodes, you can follow the next example:
- Adding "[aws_eks_node_group](#)" to your vpc.tf file in your account for substrate/infrastructure

```
resource "aws_eks_node_group" "gpu_worker" {
  cluster_name      = "adda-dev-sandbox"
  node_group_name   = "adda-dev-sandbox-managed-gpu-worker"
  node_role_arn     = module.eks-sandbox.worker_role_arn
  subnet_ids       = module.vpc_sandbox.nat_subnet_ids
  disk_size         = 250
  instance_types    = ["g5.12xlarge"]
  ami_type          = "AL2_x86_64_GPU"

  scaling_config {
    desired_size = 2
    max_size     = 4
    min_size     = 2
  }

  # taint {}

  labels = {}

  remote_access {
```

```

    ec2_ssh_key = "adda-dev-sandbox-worker"
    source_security_group_ids = [
        module.eks-sandbox.master_sg_id
    ]
}

# Ensure that IAM Role permissions are created before and deleted
# Otherwise, EKS will not be able to properly delete EC2 Instances
/*
depends_on = [
    aws_iam_role_policy_attachment.worker_node_policy,
    aws_iam_role_policy_attachment.worker_cni_policy,
    aws_iam_role_policy_attachment.worker_registry_policy,
]
*/

lifecycle {
    # create_before_destroy = true
    # prevent_destroy = true

    ignore_changes = [
        scaling_config[0].max_size,
        scaling_config[0].min_size,
        scaling_config[0].desired_size
    ]
}

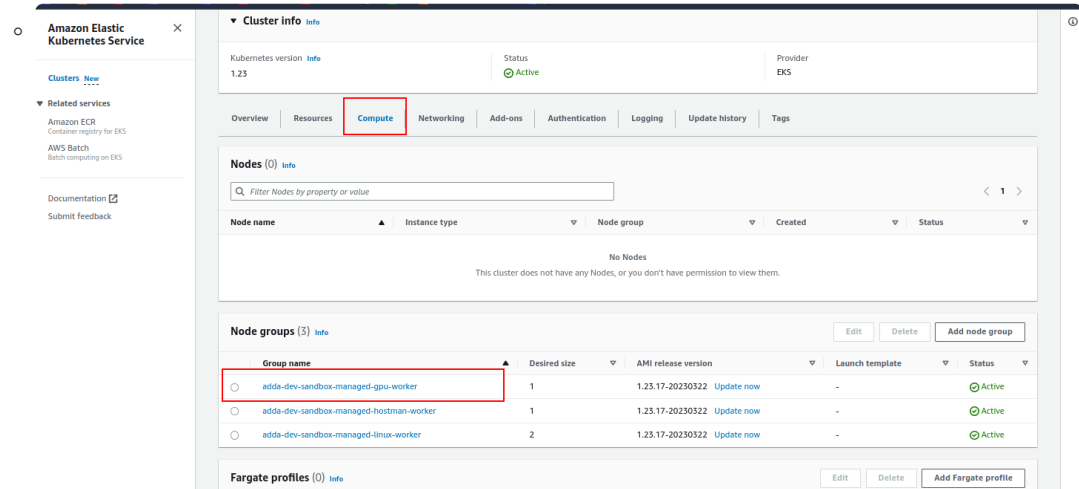
tags = {
    "epic/substrate/eks" = "adda-dev-sandbox"
}

}

```

- This procedure leaves you to enable GPU TimeSlicing sharing, as we will discuss further in this guide

- Once your terraform modification is approved, merge & applied by [#cloud-ops-support-ext](#), you'll be able to see a new worker node group attached to your EKS Compute



- To take advantage of this configuration, you require to define "nodegroup" which pod/job will be assigned as part of your NodeSelector specification as part of EKS Chart manifest

- Single pod example

```
apiVersion: v1
kind: Pod
metadata:
  name: nvidia-smi
spec:
  nodeSelector:
    eks.amazonaws.com/nodegroup: "adda-dev-sandbox-managed-gpu-worker"
  restartPolicy: OnFailure
  containers:
    - name: nvidia-smi
      image: nvidia/cuda:12.1.0-devel-ubuntu22.04
      args:
        - "nvidia-smi"
  resources:
    limits:
      nvidia.com/gpu: 2
```

- deployment/service example

-

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nvidia-smi
  labels:
    app: nvidia-smi
spec:
  replicas: 4
  nodeSelector:
    eks.amazonaws.com/nodegroup: "adda-dev-sandbox-managed-
template:
  metadata:
    name: nvidia-smi
    labels:
      app: nvidia-smi
  spec:
    containers:
      - name: nvidia-smi
        image: nvidia/cuda:12.1.0-devel-ubuntu22.04
        args:
          - "nvidia-smi"
        resources:
          limits:
            nvidia.com/gpu: 4
```

- epic-app helm example

-

```
epic-app:
  nameOverride: "hmc-canary-sandbox"
  fullnameOverride: "hmc-canary-sandbox"

  nodeSelector:
    eks.amazonaws.com/nodegroup: "adda-dev-sandbox-managed-
```

Karpenter Nodes

- Karpenter nodes are provisioned under demand, which means along your service/application is not running that computer/gpu power is not in use, the node is destroyed to avoid costs, it's more flexible if you require an increase or decrease dynamically the count of pods/replicas accessing GPU power, and removes those nodes which processing power is not being used.
- in this mode you require to modify your Karpenter configuration at ArgoCD repository, to leave substrate management systems to apply and manage EPIC systems management services, which control Karpenter configurations.
- add the [following GPU block](#) in your cluster configuration file for [ArgoCD](#)

◦

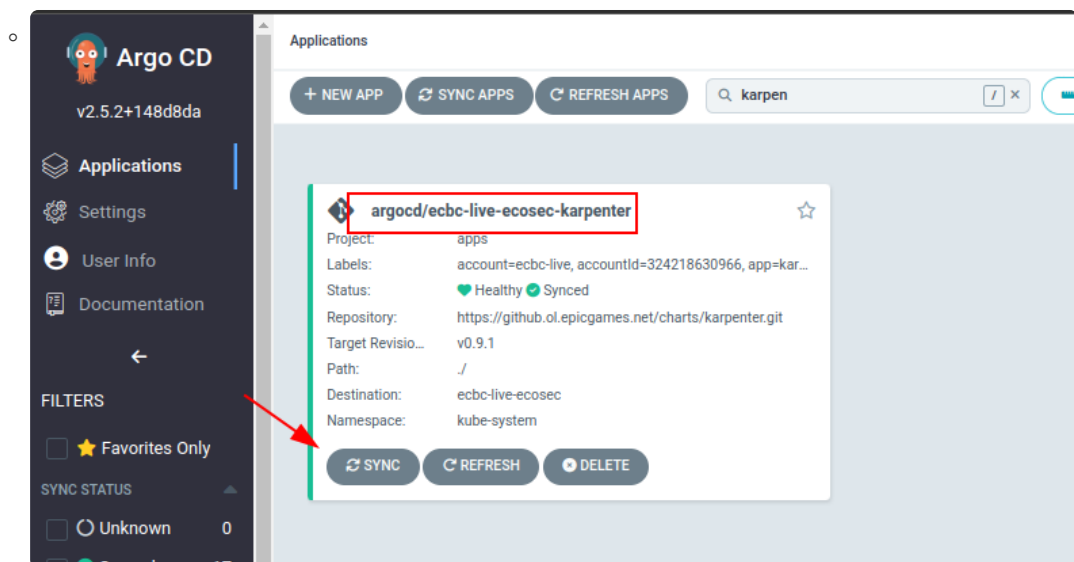
```
gpu:
  enabled: true
  labels:
    use: gpu
  providerRef:
    name: default
  requirements:
    arch:
      values:
        - amd64
    capacityType:
      values:
        - on-demand
    instanceGPUManufacturer:
      values:
        - nvidia
    instanceType:
      values:
        - g5.12xlarge
    zone:
      values: {}
  taints:
```

```

- effect: NoExecute
  key: epicgames.com/gpu
- effect: NoSchedule
  key: nvidia.com/gpu
ttlSecondsAfterEmpty: 300

```

- For more specific information about Karpenter nodes addition and management, you can review the following guides:
 - <https://confluence-epicgames.atlassian.net/wiki/spaces/IE/pages/82150715>
- Once your PR be approved & merged by [#cloud-ops-support-ext](#), you'll require to sync/apply new configuration in your cluster's Karpenter, please go to [ArgoCD GUI Page](#) and sync your Karpenter app.



- To take advantage of these configurations you require to perform the following additions to your EKS Manifest or chart, here we'll explore some examples that may adjust to your specific project requirements.
- The most important pieces of these configurations to be taken for Karpenter and provision your requested node properly and your pod/service being assigned to your provisioned node are
 - The affinity matches:

```

spec:
  affinity:

```

```

nodeAffinity:
  requiredDuringSchedulingIgnoredDuringExecution:
    nodeSelectorTerms:
      - matchExpressions:
          - key: use
            operator: In
            values:
              - gpu

```

- taint/tolerations

- ```

spec:
 ...
 tolerations:
 - effect: NoExecute
 key: epicgames.com/gpu
 operator: Exists
 - effect: NoSchedule
 key: nvidia.com/gpu
 operator: Exists

```

- Another less important but that will leave you to manage how many GPUs will be accessed by your pod/service, are the resources limits

- ```

spec:
  ...
  containers:
    ...
    resources:
      requests:
        nvidia.com/gpu: 4
      limits:
        nvidia.com/gpu: 4

```


- This is an example if you want to run a single pod (containerized application with specific arguments)

```
apiVersion: v1
kind: Pod
metadata:
  name: nvidia-smi
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: use
                operator: In
                values:
                  - gpu
  tolerations:
    - effect: NoExecute
      key: epicgames.com/gpu
      operator: Exists
    - effect: NoSchedule
      key: nvidia.com/gpu
      operator: Exists
  restartPolicy: OnFailure
  containers:
    - name: nvidia-smi
      image: nvidia/cuda:12.1.0-devel-ubuntu22.04
      args:
        - "nvidia-smi"
      resources:
        requests:
          nvidia.com/gpu: 4
        limits:
          nvidia.com/gpu: 4
```

- This is an example if you want to perform a Service Deployment (a service that requires one or more pods, long-term running application/service)
 - Here you add "replicas" value, which means how many times GPU Cores you'll need, each replica will access the count of GPU Core you specify in your resource.requests/resource.limits, (if you specify resource.requests: 2 and replica: 2, you'll require at least 4 GPU Cores to satisfy the GPU Cores demand by your Deployment)
 - Your EKS manifest should look like these pieces of code, it's probably required to adjust it depending if your service uses HELM Chart or any other kind of chart management tool.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nvidia-smi
  labels:
    app: nvidia-smi
spec:
  replicas: 4
  selector:
    matchLabels:
      app: nvidia-smi
  template:
    metadata:
      name: nvidia-smi
      labels:
        app: nvidia-smi
    spec:
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: use
```

```

        operator: In
        values:
        - gpu
    tolerations:
    - effect: NoExecute
      key: epicgames.com/gpu
      operator: Exists
    - effect: NoSchedule
      key: nvidia.com/gpu
      operator: Exists
    containers:
    - name: nvidia-smi
      image: nvidia/cuda:12.1.0-devel-ubuntu22.04
      args:
      - "nvidia-smi"
      resources:
        limits:
          nvidia.com/gpu: 4

```

- Notice that the minimum and important part to assign the job to a g5 instance is the affinity and tolerance settings:

◦

```

spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
        - matchExpressions:
          - key: use
            operator: In
            values:
            - gpu
  tolerations:
  - effect: NoExecute
    key: epicgames.com/gpu
    operator: Exists
  - effect: NoSchedule

```

```
key: nvidia.com/gpu
operator: Exists
```

- Have in mind that the Deployment manifest already specify quantity of replicas, which means how many times pod specifications will be duplicated, which means that the count of GPU Core you request to access by resources.requests in your pod will be multiplied by count of replicas specified in the Deployment manifest

◦

```
spec:
...
  replicas: 4
  template:
    spec:
      ....
      containers:
        ...
        resources:
          requests:
            nvidia.com/gpu: 4
          limits:
            nvidia.com/gpu: 4
```

- In this example you're requesting 16 GPU cores (4 replicas of 4 GPU Cores each)

◦ This is an epic-app helm chart example:

◦

```
epic-app:
  nameOverride: "hmc-canary-sandbox"
  fullnameOverride: "hmc-canary-sandbox"

  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
```

```
nodeSelectorTerms:
  - matchExpressions:
      - key: gpu
        operator: In
        values:
          - "true"
tolerations:
  - effect: NoExecute
    key: epicgames.com/gpu
    operator: Exists
  - effect: NoSchedule
    key: nvidia.com/gpu
    operator: Exists
```

Enabling GPU Support in Kubernetes

Once you have configured the options above on all the GPU nodes in your cluster, you can enable GPU support by deploying the following Daemonset:

```
$ kubectl create -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin
```

Note: This is a simple static daemonset meant to demonstrate the basic

Deployment via `helm`

The preferred method to deploy the device plugin is as a daemonset using `helm`. Instructions for installing `helm` can be found [here](#).

Begin by setting up the plugin's `helm` repository and updating it at follows:

```
$ helm repo add nvdp https://nvidia.github.io/k8s-device-plugin
$ helm repo update
```

Then verify that the latest release (`v0.14.0`) of the plugin is available:

```
$ helm search repo nvdp --devel
```

NAME	CHART VERSION	APP VERSION	DESCRIPTION
nvdp/nvidia-device-plugin	0.14.0	0.14.0	A Helm chart

Once this repo is updated, you can begin installing packages from it to deploy the `nvidia-device-plugin` helm chart.

The most basic installation command without any options is then:

```
helm upgrade -i nvdp nvdp/nvidia-device-plugin --namespace nvidia-device-p
```

Source: <https://github.com/NVIDIA/k8s-device-plugin>

Time-Slicing GPUs in Kubernetes

The NVIDIA GPU Operator enables oversubscription of GPUs through a set of extended options for the NVIDIA Kubernetes Device Plugin. GPU time-slicing enables workloads that are scheduled on oversubscribed GPUs to interleave with one another.

This mechanism for enabling time-slicing of GPUs in Kubernetes enables a system administrator to define a set of replicas for a GPU, each of which can be handed out independently to a pod to run workloads on. Unlike Multi-Instance GPU (MIG), there is no memory or fault-isolation between replicas, but for some workloads this is better than not being able to share at all. Internally, GPU time-slicing is used to multiplex workloads from replicas of the same underlying GPU.

Create your time-slicing-config.yaml file with following content

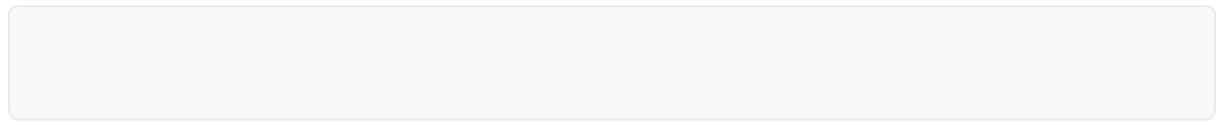
- time-slicing-config.yaml

◦

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: time-slicing-config
data:
  any: |-
    version: v1
    flags:
      migStrategy: none
    sharing:
      timeSlicing:
        renameByDefault: false
        failRequestsGreaterThanOne: false
        resources:
          - name: nvidia.com/gpu
            replicas: 4
```

- `helm upgrade -i nvdp nvdp/nvidia-device-plugin --namespace nvidia-device-plugin --create-namespace --version 0.14.0 --set-file config.map.config=time-slicing-config.yaml`

- So far we've tested this timeSlicing configuration only works properly in gpu worker-nodes method for node provisioned by karpenter it's not working yet, they're showing only the physical GPU Cores provisioned with the instance.
- For more specific information you can look at the following links
 - <https://docs.nvidia.com/datacenter/cloud-native/gpu-operator/gpu-sharing.html>
 - <https://docs.nvidia.com/datacenter/cloud-native/gpu-operator/archive/1.11.0/gpu-sharing.html>



Page Information:

Page ID: 81068307

Space: Cloud Developer Platform

Downloaded: 2025-07-12 04:07:13